



Advanced
Institute for
Artificial
Intelligence

Python Programming

<https://advancedinstitute.ai>



Pandas - Data Wrangling 101

Cleaning, Transforming & Analyzing Real-World Data

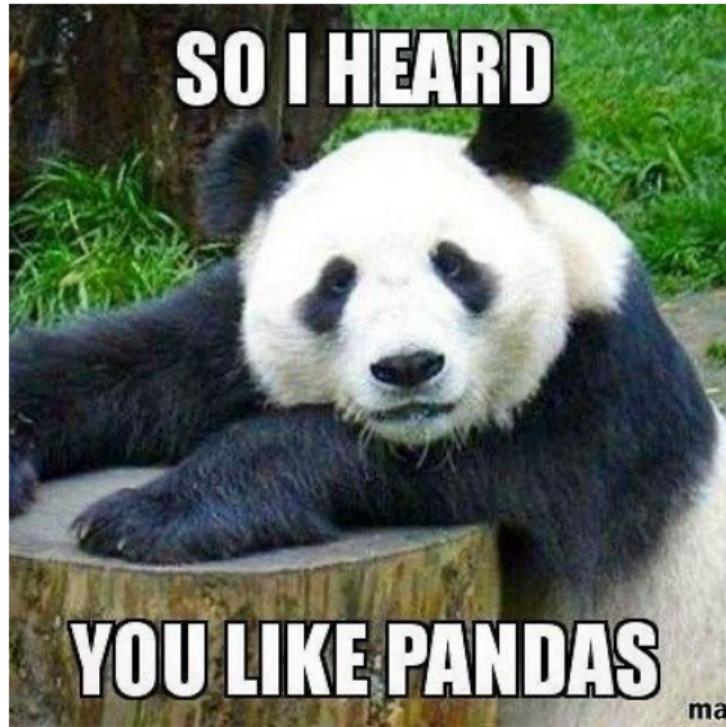
References and Image Sources

- Python for Data Analysis: Data Wrangling with Pandas, NumPy, and IPython (Book)
- Learning the Pandas Library: Python Tools for Data Munging, Analysis, and Visualization (Book)
- Official Pandas Documentation

What is Data Wrangling?

Introduction

- **Data Wrangling** (or Data Munging): The process of transforming and mapping raw data into a format suitable for analysis
- Real-world data is often:
 - Messy and incomplete
 - In the wrong format
 - Inconsistent or duplicated
 - Mixed with errors and outliers
- Data scientists spend **60-80% of their time** on data wrangling!
- Pandas is the **primary tool** for data wrangling in Python



Why Pandas?

Introduction

- **High-level** data structures for tabular and structured data
- Functions designed to make working with data **fast**, **easy**, and **expressive**
- Combines ideas from:
 - **NumPy**: high-performance array operations
 - **Spreadsheets**: intuitive data manipulation
 - **Relational databases**: SQL-like operations
- Sophisticated indexing for *reshaping*, *filtering*, and *aggregating*
- Data structures inspired by `data.frame` from R

The Story of Pandas

History and Evolution

- Created by **Wes McKinney** in 2008 while at AQR Capital Management
 - Need for high-performance, flexible tool for quantitative analysis
- First released as **open source in 2009**
- Name comes from "**Panel Data**" - econometrics term for multi-dimensional data
- Became the **de facto standard** for data manipulation in Python
- Key milestones:
 - 2011: Integration with statsmodels
 - 2015: Major refactoring and performance improvements
 - 2020: Pandas 1.0 released with stability guarantees
 - 2024: Pandas 2.x with major performance enhancements

Pandas Design Philosophy

Core Principles

- **Intuitive and easy to use** for data analysis
- Designed for **in-memory analytics** on single machines
- Prioritizes **developer productivity** over raw performance
- Built on top of **NumPy** for numerical operations
- Inspired by **R's data.frame** but with Python ergonomics
- Excellent for **exploratory data analysis (EDA)**
- Trade-off: **Convenience vs. Memory Efficiency**

Pandas and NumPy Relationship

Understanding the Foundation

- Pandas is built **on top of NumPy**
- Each Pandas column is backed by a **NumPy array**
- **NumPy**: homogeneous arrays (single data type)
 - Memory efficient, fast numerical operations
 - No index labels, position-based only
- **Pandas**: heterogeneous DataFrames (multiple data types)
 - Labeled axes (indexes), more intuitive
 - Missing data handling, alignment by labels
 - Higher-level data manipulation operations
- When to use which:
 - **NumPy**: pure numerical computations, arrays
 - **Pandas**: tabular data, mixed types, labeled data

NumPy vs Pandas - Example

Comparing Approaches

```
1 # NumPy - position-based access
2 import numpy as np
3 arr = np.array([[1, 2, 3], [4, 5, 6]])
4 arr[0, 1] # 2 (row 0, column 1)
5
6 # Pandas - label-based access
7 import pandas as pd
8 df = pd.DataFrame(arr,
9                     index=['row1', 'row2'],
10                    columns=['A', 'B', 'C'])
11 df.loc['row1', 'B'] # 2 (labeled access)
12
13 # Pandas uses NumPy underneath
14 type(df['A'].values) # <class 'numpy.ndarray'>
```

Memory Management in Pandas

How Pandas Handles Memory

- Each DataFrame column = separate **NumPy array in memory**
- **Data types matter:**
 - int64 uses 8 bytes per value
 - int32 uses 4 bytes per value
 - category dtype for repeated strings saves memory
 - object dtype (strings) most memory-intensive
- **Index overhead:** every DataFrame has an index stored separately
- **Copy-on-write (CoW):** Pandas 2.0+ reduces unnecessary copies
- Rule of thumb: **DataFrame needs 2-3x the CSV size in RAM**

Memory Optimization Example

Reducing Memory Footprint

```
1 # Check memory usage
2 df.info(memory_usage='deep')
3
4 # Before optimization
5 df['City'].dtype # object (8 bytes + string length)
6
7 # After optimization with category
8 df['City'] = df['City'].astype('category')
9 # Stores only unique values + integer codes
10
11 # Downcast numeric types
12 df['Sales'] = pd.to_numeric(df['Sales'], downcast='integer')
13
14 # Memory comparison
15 print(f"Memory saved: {original_mem - new_mem} bytes")
```



Part 1

Pandas Fundamentals

Data Structures - Series

One-Dimensional Data

- One-dimensional array with axis labels (index)
- Like a column in a spreadsheet with the **same data type**
- Index can be:
 - Numeric (default: 0, 1, 2, ...)
 - String labels
 - Dates/timestamps

apples	
0	3
1	2
2	0
3	1

Series - Basic Examples

Creating and Accessing Series

```
1 In [1]: import pandas as pd
2 In [2]: import numpy as np
3
4 # Create a Series with automatic index
5 In [3]: temperatures = pd.Series([23, 25, 22, 24])
6 0    23
7 1    25
8 2    22
9 3    24
10 dtype: int64
11
12 # Create a Series with custom index
13 In [4]: temps = pd.Series([23, 25, 22, 24],
14                           index=['Mon', 'Tue', 'Wed', 'Thu'])
15 Mon    23
16 Tue    25
17 Wed    22
```

Data Structures - DataFrame

Two-Dimensional Data

Series		Series		DataFrame	
	apples		oranges		
0	3	0	0	0	0
1	2	1	3	1	3
2	0	2	7	2	7
3	1	3	2	3	2

- ☐ Table-like structure with rows and columns
- ☐ Each column is a Series (can have different data types)
- ☐ Indexes for both rows and columns

DataFrame - Basic Example

Creating DataFrames

```
1 In [5]: data = {  
2     'City': ['São Paulo', 'Rio', 'Belo Horizonte', 'Salvador'],  
3     'Population': [12.3, 6.7, 2.5, 2.9],  
4     'Area_km2': [1521, 1200, 331, 693]  
5 }  
6 In [6]: df = pd.DataFrame(data)  
7 In [7]: df  
8          City  Population  Area_km2  
9  0      São Paulo      12.3       1521  
10 1           Rio        6.7       1200  
11 2    Belo Horizonte      2.5       331  
12 3      Salvador      2.9       693
```

Reading Data from Files

Essential Skill for Data Wrangling

□ Pandas supports multiple file formats:

- **CSV**: `pd.read_csv('data.csv')`
- **Excel**: `pd.read_excel('data.xlsx')`
- **JSON**: `pd.read_json('data.json')`
- **SQL**: `pd.read_sql(query, connection)`

□ Common parameters:

- `sep`: delimiter (default is comma for CSV)
- `header`: row number for column names
- `index_col`: column to use as row index
- `dtype`: specify data types for columns
- `parse_dates`: convert columns to datetime

Reading CSV Example

Real-World Data Loading

```
1 # Basic CSV read
2 In [8]: df = pd.read_csv('sales_data.csv')
3
4 # With custom options
5 In [9]: df = pd.read_csv('sales_data.csv',
6                         sep=';',                  # semicolon separator
7                         index_col='ID',        # use ID as index
8                         parse_dates=['Date'], # convert Date to datetime
9                         na_values=['NA', '?']) # treat as missing
10
11 # Read from URL
12 In [10]: url = 'https://example.com/data.csv'
13 In [11]: df = pd.read_csv(url)
```

Basic Data Exploration

First Look at Your Data

□ Quick inspection:

- `df.head(n)`: first n rows (default 5)
- `df.tail(n)`: last n rows
- `df.sample(n)`: random n rows

□ Data structure info:

- `df.shape`: (rows, columns)
- `df.columns`: column names
- `df.index`: row index
- `df.dtypes`: data types of each column
- `df.info()`: comprehensive overview

Exploration Example

Understanding Your Dataset

```
1 In [12]: df.head(3)
2      ID      City  Sales   Date
3  0    1  São Paulo    1500 2024-01-15
4  1    2          Rio    2300 2024-01-16
5  2    3  Salvador    1200 2024-01-17
6
7 In [13]: df.shape
8 (1250, 4)
9
10 In [14]: df.info()
11 <class 'pandas.DataFrame'>
12 RangeIndex: 1250 entries, 0 to 1249
13 Data columns (total 4 columns):
14 #   Column  Non-Null Count  Dtype  
15 ---  --       --           --      
16 0   ID      1250 non-null   int64  
17 1   City    1245 non-null   object 
```

Descriptive Statistics

- `df.describe()`: statistical summary for numeric columns
 - count, mean, std, min, 25%, 50%, 75%, max
- `df.describe(include='all')`: include non-numeric columns
- `df['column'].value_counts()`: frequency of unique values
- `df['column'].unique()`: array of unique values
- `df['column'].nunique()`: number of unique values

Accessing Data

□ Column selection:

- `df['column']`: single column (returns Series)
- `df[['col1', 'col2']]`: multiple columns (returns DataFrame)

□ Row selection by position:

- `df.iloc[0]`: first row
- `df.iloc[0:5]`: first 5 rows
- `df.iloc[[0, 2, 4]]`: specific rows

□ Row selection by label:

- `df.loc['index_label']`: row by index
- `df.loc['start':'end']`: slice by labels

Selection Examples

Practical Indexing

```
1 # Select single column
2 In [15]: df['City']
3 0      São Paulo
4 1          Rio
5 ...
6
7 # Select multiple columns
8 In [16]: df[['City', 'Sales']]
9          City   Sales
10 0    São Paulo    1500
11 1          Rio    2300
12 ...
13
14 # Select rows and columns together
15 In [17]: df.loc[0:2, ['City', 'Sales']]
16          City   Sales
17 0    São Paulo    1500
```



Part 2

Data Cleaning

Why Clean Data?

The Importance of Data Cleaning

- Real-world data is rarely clean:
 - Missing values (NaN, None, empty strings)
 - Duplicate records
 - Incorrect data types
 - Inconsistent formatting
 - Outliers and errors
- Poor data quality leads to:
 - Incorrect analysis results
 - Failed machine learning models
 - Wrong business decisions
- **Clean data is the foundation of good analysis!**

Identifying Missing Values

- ❑ Pandas represents missing data as NaN (Not a Number)

- ❑ **Detection methods:**

- `df.isnull()`: returns Boolean DataFrame
- `df.isnull().sum()`: count missing per column
- `df.notnull()`: opposite of `isnull()`

- ❑ **Visualization:**

- `df.isnull().sum().sort_values(ascending=False)`
- Shows which columns have most missing data

Missing Data Example

Finding Missing Values

```
1 In [18]: df.isnull().sum()
2 ID          0
3 City        5
4 Sales       20
5 Date         0
6 dtype: int64
7
8 # Percentage of missing values
9 In [19]: (df.isnull().sum() / len(df)) * 100
10 ID      0.00
11 City    0.40
12 Sales   1.60
13 Date    0.00
14 dtype: float64
15
16 # Rows with any missing values
17 In [20]: df[df.isnull().any(axis=1)]
```

Dealing with Missing Data

Strategies

□ Removing missing data:

- `df.dropna()`: drop rows with any NaN
- `df.dropna(axis=1)`: drop columns with any NaN
- `df.dropna(thresh=n)`: keep rows with at least n non-NaN values
- `df.dropna(subset=['col'])`: drop based on specific columns

□ Filling missing data:

- `df.fillna(value)`: fill with specific value
- `df.fillna(method='ffill')`: forward fill
- `df.fillna(method='bfill')`: backward fill
- `df.fillna(df.mean())`: fill with column mean

Filling Missing Values

Practical Examples

```
1 # Fill numeric columns with mean
2 In [21]: df['Sales'].fillna(df['Sales'].mean(), inplace=True)
3
4 # Fill categorical columns with mode
5 In [22]: df['City'].fillna(df['City'].mode()[0], inplace=True)
6
7 # Fill with specific value
8 In [23]: df['Status'].fillna('Unknown', inplace=True)
9
10 # Forward fill (use previous value)
11 In [24]: df['Price'].fillna(method='ffill', inplace=True)
12
13 # Drop rows where critical columns are missing
14 In [25]: df.dropna(subset=['ID', 'Date'], inplace=True)
```

Removing Duplicates

Finding and Removing Duplicate Records

Detection:

- `df.duplicated()`: returns Boolean Series
- `df.duplicated().sum()`: count duplicates
- `df[df.duplicated()]`: view duplicate rows

Removal:

- `df.drop_duplicates()`: remove duplicate rows
- `df.drop_duplicates(subset=['col'])`: based on specific columns
- `df.drop_duplicates(keep='first')`: keep first occurrence
- `df.drop_duplicates(keep='last')`: keep last occurrence
- `df.drop_duplicates(keep=False)`: remove all duplicates

Duplicate Handling Example

Cleaning Duplicate Data

```
1 # Check for duplicates
2 In [26]: df.duplicated().sum()
3 15
4
5 # View duplicates
6 In [27]: df[df.duplicated(keep=False)]
7      ID      City   Sales        Date
8    10      3 Salvador    1200  2024-01-17
9    25      3 Salvador    1200  2024-01-17
10
11 # Remove duplicates (keep first)
12 In [28]: df = df.drop_duplicates()
13
14 # Remove based on specific columns only
15 In [29]: df = df.drop_duplicates(subset=['ID'], keep='first')
```

Data Type Conversion

Fixing Data Types

□ Check current types: `df.dtypes`

□ Convert types:

- `df['col'].astype(int)`: convert to integer
- `df['col'].astype(float)`: convert to float
- `df['col'].astype(str)`: convert to string
- `pd.to_numeric(df['col'], errors='coerce')`: convert to numeric, NaN for errors
- `pd.to_datetime(df['col'])`: convert to datetime

□ Common issues:

- Numbers stored as strings
- Dates stored as strings
- Categories stored as objects

Type Conversion Examples

Practical Data Type Fixes

```
1 # Convert string to numeric (with error handling)
2 In [30]: df['Sales'] = pd.to_numeric(df['Sales'], errors='coerce')
3
4 # Convert to datetime
5 In [31]: df['Date'] = pd.to_datetime(df['Date'])
6
7 # Convert to category (saves memory for repeated values)
8 In [32]: df['City'] = df['City'].astype('category')
9
10 # Check types after conversion
11 In [33]: df.dtypes
12 ID           int64
13 City         category
14 Sales        float64
15 Date         datetime64[ns]
16 dtype: object
```

String Operations

Cleaning Text Data

□ Access string methods with `.str` accessor:

- `df['col'].str.lower()`: convert to lowercase
- `df['col'].str.upper()`: convert to uppercase
- `df['col'].str.strip()`: remove leading/trailing whitespace
- `df['col'].str.replace('old', 'new')`: replace text
- `df['col'].str.contains('pattern')`: check if contains text
- `df['col'].str.split('delimiter')`: split strings

□ Common use cases:

- Standardizing case
- Removing extra whitespace
- Extracting substrings



Part 3

Data Transformation

Boolean Indexing

- Use conditions to filter rows:
 - `df[df['Sales'] > 1000]`: rows where Sales > 1000
 - `df[df['City'] == 'São Paulo']`: rows for specific city
- Combine multiple conditions:
 - `&`: AND operator
 - `|`: OR operator
 - `~`: NOT operator
- **Important:** Use parentheses around each condition!
- Example: `df[(df['Sales'] > 1000) & (df['City'] == 'Rio')]`

Filtering Examples

Practical Filtering

```
1 # Simple filter
2 In [34]: high_sales = df[df['Sales'] > 2000]
3
4 # Multiple conditions (AND)
5 In [35]: result = df[(df['Sales'] > 1500) & (df['City'] == 'Rio')]
6
7 # Multiple conditions (OR)
8 In [36]: result = df[(df['City'] == 'Rio') | (df['City'] == 'São Paulo')]
9
10 # Using isin() for multiple values
11 In [37]: cities = ['Rio', 'São Paulo', 'Salvador']
12 In [38]: result = df[df['City'].isin(cities)]
13
14 # NOT condition
15 In [39]: result = df[~df['City'].isin(['Rio'])]
```

Adding Calculated Fields

□ Simple assignment:

- `df['new_col'] = value`: create with constant
- `df['total'] = df['price'] * df['quantity']`: calculation

□ Based on conditions:

- `np.where(condition, value_if_true, value_if_false)`
- Multiple conditions: `np.select()`

□ From existing columns:

- Mathematical operations
- String operations
- Date operations

Creating Columns Examples

Practical Column Creation

```
1 # Simple calculation
2 In [40]: df['Sales_k'] = df['Sales'] / 1000
3
4 # Conditional column
5 In [41]: df['Performance'] = np.where(df['Sales'] > 2000,
6                                         'High',
7                                         'Low')
8
9 # Multiple conditions
10 In [42]: conditions = [
11             df['Sales'] > 3000,
12             df['Sales'] > 2000,
13             df['Sales'] > 1000
14         ]
15 In [43]: choices = ['Excellent', 'Good', 'Average']
16 In [44]: df['Category'] = np.select(conditions, choices, default='Poor')
```

Applying Functions

Transform Data with Functions

- **apply() method:** apply function to rows or columns
 - `df['col'].apply(func)`: apply to each element in column
 - `df.apply(func, axis=0)`: apply to each column
 - `df.apply(func, axis=1)`: apply to each row
- **Lambda functions:** quick anonymous functions
 - `df['col'].apply(lambda x: x * 2)`
- **map() method:** map values using dictionary
 - `df['col'].map({'old': 'new'})`
- Use for complex transformations not possible with simple operations

Apply and Lambda Examples

Function Application

```
1 # Apply lambda to single column
2 In [45]: df['Sales_taxed'] = df['Sales'].apply(lambda x: x * 1.15)
3
4 # Apply custom function
5 In [46]: def categorize_sales(value):
6             if value > 2500:
7                 return 'High'
8             elif value > 1500:
9                 return 'Medium'
10            else:
11                return 'Low'
12
13 In [47]: df['Level'] = df['Sales'].apply(categorize_sales)
14
15 # Map values using dictionary
16 In [48]: city_codes = {'São Paulo': 'SP', 'Rio': 'RJ', 'Salvador': 'BA'}
17 In [49]: df['Code'] = df['City'].map(city_codes)
```

Ordering Your Data

□ Sort by values:

- `df.sort_values('col')`: sort by single column
- `df.sort_values(['col1', 'col2'])`: sort by multiple columns
- `ascending=False`: descending order

□ Sort by index:

- `df.sort_index()`: sort by row index

□ Parameters:

- `inplace=True`: modify original DataFrame
- `na_position`: where to put NaN values ('first' or 'last')

Sorting Examples

Practical Sorting

```
1 # Sort by single column (ascending)
2 In [50]: df_sorted = df.sort_values('Sales')
3
4 # Sort descending
5 In [51]: df_sorted = df.sort_values('Sales', ascending=False)
6
7 # Sort by multiple columns
8 In [52]: df_sorted = df.sort_values(['City', 'Sales'],
9                                ascending=[True, False])
10
11 # Sort and modify original
12 In [53]: df.sort_values('Date', inplace=True)
13
14 # Reset index after sorting
15 In [54]: df = df.sort_values('Sales').reset_index(drop=True)
```

Renaming Columns

Better Column Names

Rename specific columns:

- `df.rename(columns={'old': 'new'})`

Rename all columns:

- `df.columns = ['name1', 'name2', ...]`

Clean column names:

- Remove spaces: `df.columns.str.replace(' ', '_')`
- Lowercase: `df.columns.str.lower()`
- Strip whitespace: `df.columns.str.strip()`

Good practice: use lowercase with underscores



Part 4

Aggregation & Grouping

Split-Apply-Combine Pattern

- ❑ **GroupBy** is one of the most powerful Pandas features
- ❑ Three-step process:
 - **Split**: divide data into groups based on criteria
 - **Apply**: apply function to each group independently
 - **Combine**: combine results into a data structure
- ❑ Basic syntax: `df.groupby('column')`
- ❑ Similar to SQL's GROUP BY

GroupBy Basics

Simple Aggregations

```
1 # Group by single column and calculate mean
2 In [55]: df.groupby('City')['Sales'].mean()
3 City
4 Belo Horizonte    1850.5
5 Rio               2150.3
6 Salvador          1320.8
7 São Paulo         2850.2
8 Name: Sales, dtype: float64
9
10 # Group by and get multiple statistics
11 In [56]: df.groupby('City')['Sales'].agg(['mean', 'sum', 'count'])
12
13           mean      sum   count
14 City
15 Belo Horizonte  1850.5  22206      12
16 Rio              2150.3  64509      30
17 ...
```

Aggregation Functions

Common Aggregations

□ After groupby, apply aggregation functions:

- `.sum()`: sum of values
- `.mean()`: average
- `.median()`: median value
- `.min()`, `.max()`: minimum/maximum
- `.count()`: number of non-null values
- `.std()`, `.var()`: standard deviation/variance

□ **agg() method:** apply multiple functions at once

- `df.groupby('City').agg(['mean', 'sum', 'count'])`

Multiple Aggregations

Advanced Grouping

```
1 # Different aggregations for different columns
2 In [57]: df.groupby('City').agg({
3     'Sales': ['mean', 'sum'],
4     'Quantity': 'count',
5     'Profit': ['min', 'max']
6 })
7
8 # Group by multiple columns
9 In [58]: df.groupby(['City', 'Product'])['Sales'].sum()
10
11 # Custom aggregation function
12 In [59]: df.groupby('City')['Sales'].agg(
13     lambda x: x.max() - x.min()
14 )
```

Excel-Style Pivot Tables

- **pivot_table()**: create spreadsheet-style pivot tables
- Key parameters:
 - values: column to aggregate
 - index: column(s) for rows
 - columns: column(s) for columns
 - aggfunc: aggregation function (default: mean)
 - fill_value: value for missing combinations
- Useful for cross-tabulation and multi-dimensional aggregation

Pivot Table Example

Creating Pivot Tables

```
1 # Simple pivot table
2 In [60]: pd.pivot_table(df,
3                         values='Sales',
4                         index='City',
5                         columns='Month',
6                         aggfunc='sum')
7 Month          Jan      Feb      Mar
8 City
9 Rio           12300.0  15200.0  14100.0
10 Salvador     8900.0   9200.0  10500.0
11 São Paulo    18500.0  19200.0  21300.0
12
13 # With multiple aggregations
14 In [61]: pd.pivot_table(df, values='Sales',
15                         index='City',
16                         aggfunc=['sum', 'mean', 'count'])
```



Part 5

Combining Datasets

Concatenating DataFrames

Stacking Data

- **concat()**: combine DataFrames along an axis
- Vertical concatenation (stack rows):
 - `pd.concat([df1, df2])`
 - Use when combining data with same columns
- Horizontal concatenation (stack columns):
 - `pd.concat([df1, df2], axis=1)`
- Parameters:
 - `ignore_index=True`: create new sequential index
 - `keys`: create hierarchical index

Concat Examples

Combining DataFrames

```
1 # Vertical concatenation (adding rows)
2 In [62]: jan_data = pd.read_csv('january.csv')
3 In [63]: feb_data = pd.read_csv('february.csv')
4 In [64]: combined = pd.concat([jan_data, feb_data], ignore_index=True)
5
6 # Horizontal concatenation (adding columns)
7 In [65]: sales = pd.DataFrame({'Sales': [100, 200]})
8 In [66]: costs = pd.DataFrame({'Costs': [60, 120]})
9 In [67]: combined = pd.concat([sales, costs], axis=1)
10
11      Sales  Costs
12      0      100     60
13      1      200     120
```

SQL-Style Joins

- **merge():** combine DataFrames based on common columns
- Types of joins:
 - `how='inner'`: keep only matching rows (default)
 - `how='left'`: keep all rows from left DataFrame
 - `how='right'`: keep all rows from right DataFrame
 - `how='outer'`: keep all rows from both
- Parameters:
 - `on`: column name(s) to join on
 - `left_on`, `right_on`: different column names

Merge Examples

Joining DataFrames

```
1 # Inner join (default)
2 In [68]: customers = pd.DataFrame({'ID': [1,2,3], 'Name': ['Ana','Bob','Carol']})
3 In [69]: orders = pd.DataFrame({'ID': [1,2,4], 'Amount': [100,200,150]})
4 In [70]: pd.merge(customers, orders, on='ID')
5      ID    Name   Amount
6      0      Ana     100
7      1      Bob     200
8
9 # Left join (keep all customers)
10 In [71]: pd.merge(customers, orders, on='ID', how='left')
11      ID    Name   Amount
12     0      Ana    100.0
13     1      Bob    200.0
14     2      Carol   NaN
```

When to Use Each Method

Concat vs Merge

Use `concat()` when:

- Combining data with same structure (same columns)
- Stacking data vertically (e.g., monthly files)
- Adding new columns to existing data

Use `merge()` when:

- Joining related data from different sources
- Combining based on a key/relationship
- Need SQL-like join operations

Think: concat for simple stacking, merge for relational joins



Wrap-Up

Performance and Parallel Processing

Limitations for Large Datasets

- **In-memory only:** entire dataset must fit in RAM
 - Typical limit: **5-10 GB** on consumer hardware
 - With 32 GB RAM: safely handle **10-15 GB datasets**
- **Single-threaded by default:** doesn't utilize all CPU cores
- **Memory copies:** operations often create intermediate copies
- **String overhead:** object dtype stores Python objects (expensive)
- **Index overhead:** can be 20-30% of total memory
- When Pandas struggles:
 - Datasets larger than available RAM
 - Operations requiring multiple passes over data
 - Need for distributed processing

When You Need More Power

For larger-than-memory datasets:

- **Dask**: parallel computing with Pandas-like API
- **Vaex**: lazy evaluation, out-of-core processing
- **Polars**: Rust-based, faster and more memory efficient

For distributed computing:

- **PySpark**: distributed DataFrames across cluster
- **Ray**: distributed Python for large-scale data

For columnar analytics:

- **Apache Arrow**: in-memory columnar format
- **DuckDB**: SQL on Pandas DataFrames

Alternative: Dask

Parallel Pandas

- Extends Pandas to **parallel and distributed computing**
- **Familiar API:** looks like Pandas, scales like Spark
- Use cases:
 - Datasets **larger than RAM** (10s to 100s of GB)
 - Parallel processing on **multi-core machines**
 - Distributed processing on **clusters**
- **Lazy evaluation:** computes only when needed
- Limitations:
 - Not all Pandas operations supported
 - Overhead for small datasets
 - Requires understanding of parallel computing

Dask Example

Using Dask for Large Datasets

```
1 import dask.dataframe as dd
2
3 # Read large CSV that doesn't fit in memory
4 df = dd.read_csv('huge_file.csv')
5
6 # Dask API similar to Pandas
7 result = df.groupby('City')['Sales'].mean()
8
9 # Lazy evaluation - nothing computed yet!
10 # Compute triggers execution
11 result.compute()
12
13 # Parallel processing across multiple cores
14 df = dd.read_csv('data/*.csv') # Read many files
15 df.groupby('Category').sum().compute()
```

Alternative: Vaex

Out-of-Core DataFrames

- Designed for **lazy, out-of-core DataFrames**
- **Memory-mapped files:** doesn't load entire dataset
- Key features:
 - **Instant visualization** of billion-row datasets
 - **Zero memory copy** operations
 - **Expression system** - only computes what's needed
 - Excellent for **exploratory analysis** of huge datasets
- Best for:
 - Numerical data (astronomy, finance, sensor data)
 - Aggregations and filtering on large datasets
 - When you need speed without loading everything
- Trade-offs:
 - Smaller ecosystem than Pandas
 - Works best with specific file formats (**HDF5, Apache Arrow**)

Vaex Example

Billion Row Processing

```
1 import vaex
2
3 # Open large file - loads metadata only, not data
4 df = vaex.open('huge_dataset.hdf5')
5
6 # Instant statistics on billion rows
7 print(df.mean('Sales')) # No data loaded!
8
9 # Lazy filtering and aggregation
10 filtered = df[df.Sales > 1000]
11 result = filtered.groupby('City').agg({'Sales': 'mean'})
12
13 # Only computes when exporting
14 result.export('result.csv')
15
16 # Memory efficient - processes in chunks
17 print(df.memory_usage()) # Minimal!
```

Alternative: Polars

Next-Generation DataFrame Library

- Written in **Rust**, Python bindings
- **Multi-threaded by default** - uses all CPU cores
- Key advantages:
 - **2-10x faster** than Pandas for most operations
 - **Lower memory usage** through better algorithms
 - **Lazy evaluation** with query optimization
 - **Type system** prevents common errors
- API differences from Pandas:
 - Method chaining preferred over in-place operations
 - Expressions instead of direct column access
- Best for:
 - New projects where performance matters
 - Data that fits in memory but needs speed
 - ETL pipelines and data transformations

Polars Example

High-Performance Analytics

```
1 import polars as pl
2
3 # Read CSV - uses all CPU cores
4 df = pl.read_csv('sales_data.csv')
5
6 # Lazy evaluation with query optimization
7 result = (
8     df.lazy()
9     .filter(pl.col('Sales') > 1000)
10    .groupby('City')
11    .agg([
12        pl.col('Sales').mean().alias('avg_sales'),
13        pl.col('Sales').sum().alias('total_sales')
14    ])
15    .collect() # Execute optimized query
16 )
17
```

Comparison: When to Use What?

Decision Guide

Pandas:

- Data **fits in memory** (< 10 GB)
- Exploratory data analysis, prototyping
- Maximum ecosystem support and documentation

Polars:

- Need **speed** and **efficiency** on single machine
- Data fits in memory but performance matters
- New projects, willing to learn new API

Dask:

- Data **larger than RAM** (10-1000 GB)
- Need to scale to multiple machines
- Want Pandas-like API

Comparison: When to Use What? (cont.)

Decision Guide

Vaex:

- **Billions of rows**, numerical data
- Exploratory analysis without loading full dataset
- Streaming data processing

PySpark:

- **Truly big data** (TBs to PBs)
- Existing Spark infrastructure
- Complex distributed workflows

Rule of thumb:

- Start with Pandas
- Optimize with Polars when needed
- Scale with Dask/Spark when data grows

Performance Tips for Pandas

Making Pandas Faster

Use appropriate data types:

- category for repeated strings
- Downcast integers (int64 → int32 → int16)

Vectorized operations: avoid loops

- Use built-in methods instead of apply()
- df['col'] * 2 instead of df['col'].apply(lambda x: x * 2)

Read data efficiently:

- Specify dtypes when reading CSV
- Use usecols to read only needed columns
- Consider Parquet format instead of CSV

Use Copy-on-Write (CoW): Pandas 2.0+

Chunk large operations: process in batches

Performance Optimization Example

Before and After

```
1 # SLOW - Python loop
2 for idx, row in df.iterrows():
3     df.loc[idx, 'result'] = row['a'] * row['b']
4
5 # FAST - Vectorized operation
6 df['result'] = df['a'] * df['b']
7
8 # SLOW - apply with lambda
9 df['price_taxed'] = df['price'].apply(lambda x: x * 1.15)
10
11 # FAST - Direct vectorization
12 df['price_taxed'] = df['price'] * 1.15
13
14 # Memory efficient reading
15 df = pd.read_csv('data.csv',
16                  dtype={'City': 'category'},
17                  usecols=['City', 'Sales'],
```

The Future of Pandas

What's Next

Pandas 2.x and beyond:

- Copy-on-Write (CoW) as default
- Better memory management
- Improved performance through refactoring
- PyArrow backend for better efficiency

Integration with Arrow ecosystem

Continued focus on:

- User experience and API consistency
- Better error messages
- Type hints and modern Python features

Pandas remains **essential** even as alternatives emerge

- Lingua franca of data analysis in Python
- Gateway to more specialized tools

Key Takeaways - Architecture & Scaling

Summary

- ❑ Pandas is built on NumPy, trades memory for convenience
- ❑ Designed for in-memory, single-machine analytics
- ❑ Memory footprint: **2-3x the data size**
- ❑ Practical limit: **5-15 GB datasets** on typical hardware
- ❑ For larger data:
 - Polars: speed on single machine
 - Dask: larger-than-memory, multi-machine
 - Vaex: billion-row exploration
 - PySpark: truly big data (TB+)
- ❑ **Pandas is not going away** - it's the foundation!
- ❑ Learn Pandas first, then explore alternatives as needed



Wrap-Up

Key Takeaways & Best Practices

Data Wrangling Workflow

Typical Process

- **1. Load Data:** read from files or databases
- **2. Explore:** understand structure, types, and issues
 - `head()`, `info()`, `describe()`, `shape`, `dtypes`
- **3. Clean:** handle missing data, duplicates, types
 - `dropna()`, `fillna()`, `drop_duplicates()`, `astype()`
- **4. Transform:** filter, create columns, apply functions
 - Boolean indexing, `apply()`, new columns
- **5. Aggregate:** group and summarize
 - `groupby()`, `pivot_table()`
- **6. Combine:** merge multiple datasets
 - `concat()`, `merge()`

Tips for Effective Data Wrangling

- **Always explore before cleaning**
 - Understand your data first
- **Document your assumptions**
 - How you handle missing data, why you remove outliers
- **Keep original data intact**
 - Work on copies: `df_clean = df.copy()`
- **Use meaningful variable names**
- **Chain operations carefully**
 - Method chaining is powerful but can be hard to debug
- **Save intermediate results**
 - Especially after time-consuming operations

Common Mistakes to Avoid

Pitfalls

- Not using `inplace` parameter correctly**
 - Remember: most operations return a new DataFrame
- Forgetting to reset index after sorting/filtering**
 - Use `reset_index(drop=True)`
- Not handling missing data properly**
 - Understand why data is missing before deciding what to do
- Ignoring data types**
 - Wrong types lead to wrong results
- Not checking for duplicates**
- Overwriting original data without backup**

Resources for Learning More

Continue Your Journey

□ Official Documentation:

- pandas.pydata.org/docs

□ Books:

- Python for Data Analysis by Wes McKinney (Pandas creator)

□ Practice Datasets:

- Kaggle datasets: kaggle.com/datasets
- UCI Machine Learning Repository

□ Interactive Learning:

- DataCamp, Coursera, edX courses on Pandas

□ Practice is key! Work with real datasets

Questions?